

# Neural Networks & Deep Learning: ICP8

Name: Lalitha Sowjanya Kamuju

ID: 700747213

```
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import fashion_mnist
import numpy as np

# this is the size of our encoded representations
encoding_dim = 32 # 32 floats -> compression factor 24.5, assuming the input is 784 floats

# this is our input placeholder
input_img = Input(shape=(784,))
# "encoded" is the encoded representation of the input
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)
# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])

(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

autoencoder.fit(x_train, x_train,
                epochs=5,
                batch_size=256,
                shuffle=True,
                validation_data=(x_test, x_test))
```

```

Epoch 1/5
235/235 [=====] - 4s 15ms/step - loss: 0.6954 - accuracy: 9.5000e-04 - val_loss: 0.6952 - val_accuracy: 0.0010
Epoch 2/5
235/235 [=====] - 4s 15ms/step - loss: 0.6951 - accuracy: 9.6667e-04 - val_loss: 0.6949 - val_accuracy: 0.0011
Epoch 3/5
235/235 [=====] - 3s 13ms/step - loss: 0.6948 - accuracy: 9.5000e-04 - val_loss: 0.6946 - val_accuracy: 0.0011
Epoch 4/5
235/235 [=====] - 3s 11ms/step - loss: 0.6945 - accuracy: 9.5000e-04 - val_loss: 0.6944 - val_accuracy: 0.0012
Epoch 5/5
235/235 [=====] - 3s 14ms/step - loss: 0.6942 - accuracy: 8.8333e-04 - val_loss: 0.6941 - val_accuracy: 0.0012

<keras.src.callbacks.History at 0x7977d7fb4d90>

```

```

from keras.layers import Input, Dense
from keras.models import Model

# This is the size of our encoded representation
encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

# This is our input placeholder
input_img = Input(shape=(784,))

# "encoded" is the encoded representation of the input
encoded1 = Dense(128, activation='relu')(input_img)
encoded2 = Dense(encoding_dim, activation='relu')(encoded1)

# "decoded" is the lossy reconstruction of the input
decoded1 = Dense(128, activation='relu')(encoded2)
decoded2 = Dense(784, activation='sigmoid')(decoded1)

# This model maps an input to its reconstruction
autoencoder = Model(input_img, decoded2)

# This model maps an input to its encoded representation
encoder = Model(input_img, encoded2)

# This is our decoder model
encoded_input = Input(shape=(encoding_dim,))
decoder_layer1 = autoencoder.layers[-2]
decoder_layer2 = autoencoder.layers[-1]
decoder = Model(encoded_input, decoder_layer2(decoder_layer1(encoded_input)))

# Compile the model
autoencoder.compile(optimizer='adadelata', loss='binary_crossentropy', metrics=['accuracy'])

# Load the MNIST dataset
from keras.datasets import mnist, fashion_mnist
import numpy as np
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

```

```

import numpy as np
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

# Normalize and flatten the data
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

# Train the autoencoder
autoencoder.fit(x_train, x_train,
               epochs=5,
               batch_size=256,
               shuffle=True,
               validation_data=(x_test, x_test))

```

```

Epoch 1/5
235/235 [=====] - 6s 21ms/step - loss: 0.6931 - accuracy: 0.0022 - val_loss: 0.6931 - val_accuracy: 0.0017
Epoch 2/5
235/235 [=====] - 5s 19ms/step - loss: 0.6931 - accuracy: 0.0022 - val_loss: 0.6930 - val_accuracy: 0.0017
Epoch 3/5
235/235 [=====] - 4s 17ms/step - loss: 0.6930 - accuracy: 0.0022 - val_loss: 0.6930 - val_accuracy: 0.0017
Epoch 4/5
235/235 [=====] - 5s 21ms/step - loss: 0.6930 - accuracy: 0.0022 - val_loss: 0.6929 - val_accuracy: 0.0017
Epoch 5/5
235/235 [=====] - 5s 21ms/step - loss: 0.6929 - accuracy: 0.0022 - val_loss: 0.6929 - val_accuracy: 0.0017

<keras.src.callbacks.History at 0x7977e2ccf790>

```

```
import matplotlib.pyplot as plt

# Get the reconstructed images for the test set
reconstructed_imgs = autoencoder.predict(x_test)


# Choose a random image from the test set
n = 10 # index of the image to be plotted
plt.figure(figsize=(10, 5))

# Plot the original image
ax = plt.subplot(1, 2, 1)
plt.imshow(x_test[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Original Image")


# Plot the reconstructed image
ax = plt.subplot(1, 2, 2)
plt.imshow(reconstructed_imgs[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Reconstructed Image")

plt.show()
```

Original Image



Reconstructed Image



```

from keras.layers import Input, Dense
from keras.models import Model

# this is the size of our encoded representations
encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

# this is our input placeholder
input_img = Input(shape=(784,))
# "encoded" is the encoded representation of the input
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)
# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])
from keras.datasets import fashion_mnist
import numpy as np
(x_train, _), (x_test, _) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

#introducing noise
noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

autoencoder.fit(x_train_noisy, x_train,
                epochs=10,
                batch_size=256,
                shuffle=True,
                validation_data=(x_test_noisy, x_test_noisy))

```

```
Epoch 1/10
235/235 [=====] - 5s 16ms/step - loss: 0.6971 - accuracy: 0.0011 - val_loss: 0.6970 - val_accuracy: 0.0014
Epoch 2/10
235/235 [=====] - 3s 12ms/step - loss: 0.6968 - accuracy: 0.0010 - val_loss: 0.6967 - val_accuracy: 0.0015
Epoch 3/10
235/235 [=====] - 4s 17ms/step - loss: 0.6964 - accuracy: 0.0011 - val_loss: 0.6964 - val_accuracy: 0.0015
Epoch 4/10
235/235 [=====] - 6s 26ms/step - loss: 0.6961 - accuracy: 0.0011 - val_loss: 0.6960 - val_accuracy: 0.0015
Epoch 5/10
235/235 [=====] - 3s 13ms/step - loss: 0.6958 - accuracy: 0.0011 - val_loss: 0.6957 - val_accuracy: 0.0015
Epoch 6/10
235/235 [=====] - 3s 11ms/step - loss: 0.6955 - accuracy: 0.0011 - val_loss: 0.6954 - val_accuracy: 0.0014
Epoch 7/10
235/235 [=====] - 3s 11ms/step - loss: 0.6952 - accuracy: 0.0012 - val_loss: 0.6951 - val_accuracy: 0.0014
Epoch 8/10
235/235 [=====] - 4s 17ms/step - loss: 0.6949 - accuracy: 0.0011 - val_loss: 0.6948 - val_accuracy: 0.0016
Epoch 9/10
235/235 [=====] - 3s 12ms/step - loss: 0.6946 - accuracy: 0.0011 - val_loss: 0.6946 - val_accuracy: 0.0016
Epoch 10/10
235/235 [=====] - 3s 11ms/step - loss: 0.6943 - accuracy: 0.0012 - val_loss: 0.6943 - val_accuracy: 0.0017

<keras.src.callbacks.History at 0x7977d71fb820>
```

```
import matplotlib.pyplot as plt

# Get the reconstructed images for the test set
reconstructed_imgs = autoencoder.predict(x_test_noisy)


# Choose a random image from the test set
n = 10 # index of the image to be plotted
plt.figure(figsize=(10, 5))

# Plot the original noisy image
ax = plt.subplot(1, 2, 1)
plt.imshow(x_test_noisy[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Noisy Image")


# Plot the reconstructed image
ax = plt.subplot(1, 2, 2)
plt.imshow(reconstructed_imgs[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Reconstructed Image")

plt.show()
```

Noisy Image



Reconstructed Image



```
import matplotlib.pyplot as plt


# Train the autoencoder
history = autoencoder.fit(x_train_noisy, x_train,
                           epochs=10,
                           batch_size=256,
                           shuffle=True,
                           validation_data=(x_test_noisy, x_test_noisy))


# Plot the loss
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='test')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend()
plt.show()

# Plot the accuracy
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='test')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```



```
Epoch 1/10
235/235 [=====] - 5s 22ms/step - loss: 0.6941 - accuracy: 0.0011 - val_loss: 0.6940 - val_accuracy: 0.0017
Epoch 2/10
235/235 [=====] - 3s 12ms/step - loss: 0.6938 - accuracy: 0.0012 - val_loss: 0.6937 - val_accuracy: 0.0017
Epoch 3/10
235/235 [=====] - 3s 13ms/step - loss: 0.6935 - accuracy: 0.0012 - val_loss: 0.6935 - val_accuracy: 0.0017
Epoch 4/10
235/235 [=====] - 3s 13ms/step - loss: 0.6933 - accuracy: 0.0012 - val_loss: 0.6932 - val_accuracy: 0.0018
Epoch 5/10
235/235 [=====] - 4s 17ms/step - loss: 0.6930 - accuracy: 0.0012 - val_loss: 0.6930 - val_accuracy: 0.0018
Epoch 6/10
235/235 [=====] - 3s 12ms/step - loss: 0.6928 - accuracy: 0.0012 - val_loss: 0.6927 - val_accuracy: 0.0018
Epoch 7/10
235/235 [=====] - 3s 12ms/step - loss: 0.6925 - accuracy: 0.0012 - val_loss: 0.6925 - val_accuracy: 0.0018
Epoch 8/10
235/235 [=====] - 3s 12ms/step - loss: 0.6923 - accuracy: 0.0013 - val_loss: 0.6922 - val_accuracy: 0.0018
Epoch 9/10
235/235 [=====] - 4s 16ms/step - loss: 0.6920 - accuracy: 0.0013 - val_loss: 0.6920 - val_accuracy: 0.0017
Epoch 10/10
235/235 [=====] - 3s 12ms/step - loss: 0.6918 - accuracy: 0.0013 - val_loss: 0.6917 - val_accuracy: 0.0017
```





Git Link : <https://github.com/sowjanya-kamuju/Assignment8>

Video Link : <https://vimeo.com/925696153/c49e0ee7c0?share=copy>